

AKSHAY DUBEY

Mumbai, India

☎ +91-8779197850 ✉ akshaydubey1818@gmail.com 💼 [akshay-dubey-2060](#) 🌐 [akshay-dubey](#)

Summary

Full Stack Engineer building production LLM systems. Currently at **Fynd** (Reliance-backed), developing multi-model AI orchestration, RAG pipelines with semantic vector search, and deterministic scoring engines for an AI-powered hiring platform. 4+ years building scalable applications with TypeScript, Node.js, and PostgreSQL.

Skills

Languages: TypeScript, Python, Node.js, JavaScript

AI/LLM: OpenAI API, OpenRouter, LangChain, LangGraph, Multi-Model Orchestration, RAG Pipelines, Embeddings

AI Infrastructure: Langfuse (LLM Observability), MCP Protocol, Circuit Breakers, PG Boss Queues, Batch Processing

Databases: PostgreSQL, pgvector, Redis, MongoDB, Prisma ORM, Semantic Search

Backend: Express.js, GraphQL, REST APIs, JWT Authentication

Cloud & DevOps: GCP (Cloud Run), Docker, CI/CD Pipelines, Vercel, Supabase

Experience

Fynd (Shopsense Retail Technologies)

Mumbai, India

AI Full Stack Engineer

March 2025 – Present

- Implemented **multi-model LLM council system** using OpenRouter coordinating GPT-4, Claude, and Gemini through a 3-stage pipeline: parallel independent responses, peer ranking with critique, and chairman synthesis for consensus-based hiring decisions.
- Built **multi-vector RAG pipeline** using OpenAI text-embedding-3-large (3072 dimensions) with 7 semantic vectors (skills, role, domain, seniority, functional, education, full-profile) stored in **pgvector** for candidate-job matching with weighted similarity scoring.
- Developed **9-factor deterministic scoring engine** with configurable weighted factors, cap rules, and modifiers — aggregating AI-extracted signals into explainable match scores with quality gates for data confidence and JD completeness.
- Contributed to **asynchronous processing pipeline** using PG Boss queue with multi-stage orchestration (Parse, Deduplicate, Score), exponential backoff retry, and singleton deduplication for bulk candidate processing.
- Integrated **Langfuse for LLM observability** — enabling trace debugging across multi-model council sessions, token usage monitoring, and AI pipeline performance analysis.
- Implemented tool-enabled AI agents using **MCP (Model Context Protocol)** with circuit breakers on AI/storage services and batch processing (8 items per AI call) for resilient document processing.
- Built **hybrid pool scoring system** combining per-dimension cosine similarity scores with structured candidate signals — enabling zero-AI-call scoring at search time using pre-computed embeddings and index signals.

Osian Infotech Pvt Ltd

Mumbai, India

Frontend Developer

March 2022 – February 2025

- Built and maintained production **React/TypeScript** applications, delivering responsive interfaces optimized for performance across devices.
- Optimized application performance using React memoization, lazy loading, and code splitting, achieving improved **Core Web Vitals** scores.
- Implemented **Redux and RTK Query** for state management and API caching, ensuring consistent data flow across components.

Projects

Content Creation FYI | *Next.js, PostgreSQL, Supabase, Prisma, NextAuth, TypeScript*

[contentcreation.fyi](#)

- Built full-stack **Next.js** platform with PostgreSQL on Supabase, Prisma ORM, and API routes serving 200+ curated tools for content creators.
- Implemented **NextAuth** authentication with JWT sessions, protected API routes, and engagement features including bookmarking, upvoting, and server-side rendering for SEO optimization.

Education

Mumbai University

Maharashtra, India

Bachelor of Science in Computer Science (BSc.CS); CGPA: 7.27

July 2019 – March 2022